

# A Common Machine-Readable Vocabulary and Knowledge Base Supporting Multiple Programming Models

March 16th, 2021

ASCR Workshop on Reimagining Codesign



Chunhua “Leo” Liao



LLNL-PRES-820348

This work was performed under the auspices of the U.S. Department of Energy by Lawrence Livermore National Laboratory under contract DE-AC52-07NA27344. Lawrence Livermore National Security, LLC


# The Need

## Programming Models

- Goal: providing high-level abstractions for programmers to express algorithms+ generating correct, fast and efficient executables
- Components: languages + compilers + runtime systems + performance and correctness tools ...

A large body of knowledge is required to build efficient programming models

- Applications: domain-specific optimization knowledge, performance models
- Programming Language: semantics for language constructs
- Compilers: eligibility and profitability of optimizations
- Libraries: standard and user-defined semantics, such as read-only, aliasing, continuous memory
- Hardware: architectural features and corresponding optimization opportunities




# The State of The Art

Knowledge needed for programming models is implicitly assumed or represented using ad-hoc approaches

- Vocabularies: different communities may talk about the same concepts, using very different terms. How to collaborate then?
- Informal: not accessible, not interoperable, no verifiable, not scalable, ....
- Information gets lost across different layers
- Not reusable across different programming models

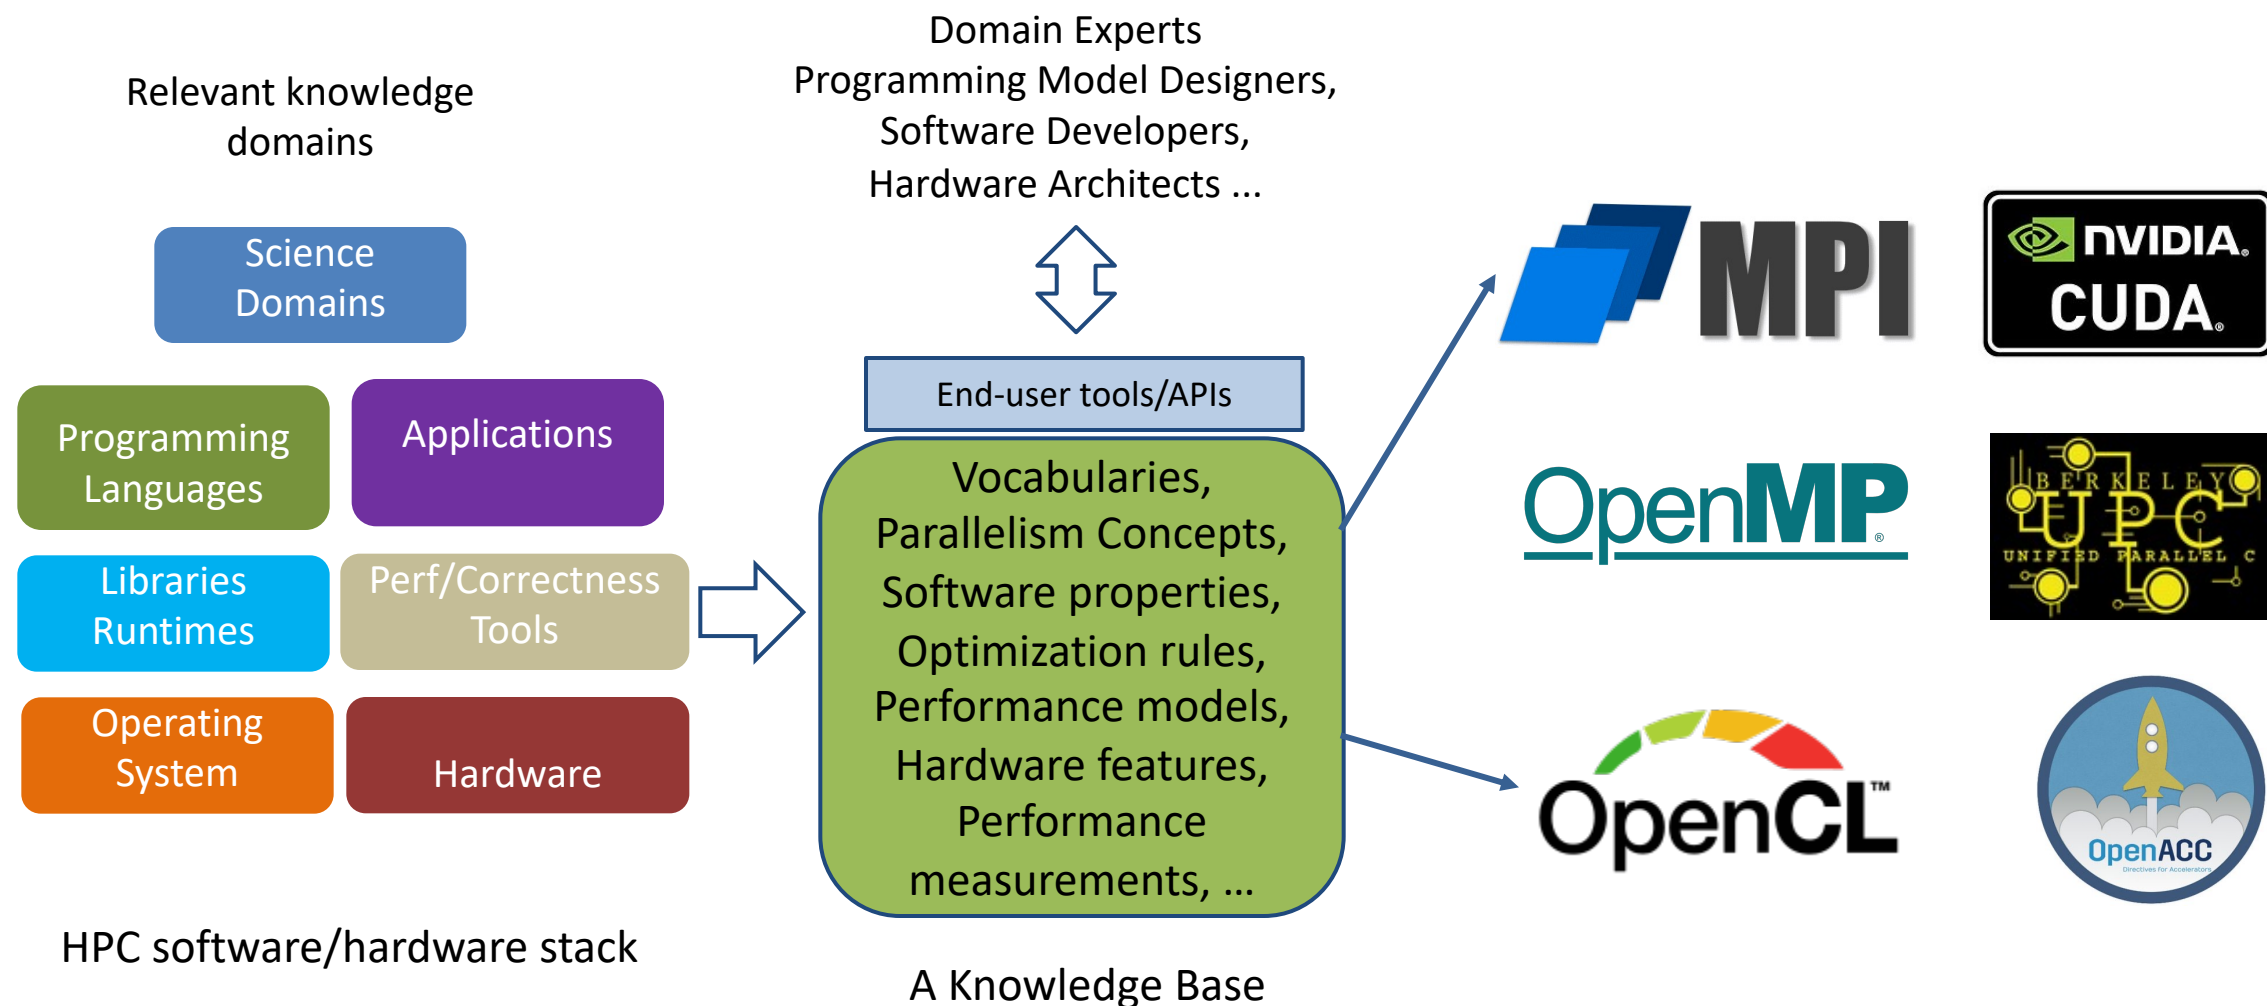
Every programming model, for every domain

- Re-inventing the wheels to bridge the semantics/knowledge gap



HPC software/hardware stack

# Reimagining the Ecosystem of Programming Models





# Solution: Knowledge Graphs= Vocabularies + Ontologies+...

Knowledge can be modeled using knowledge graphs

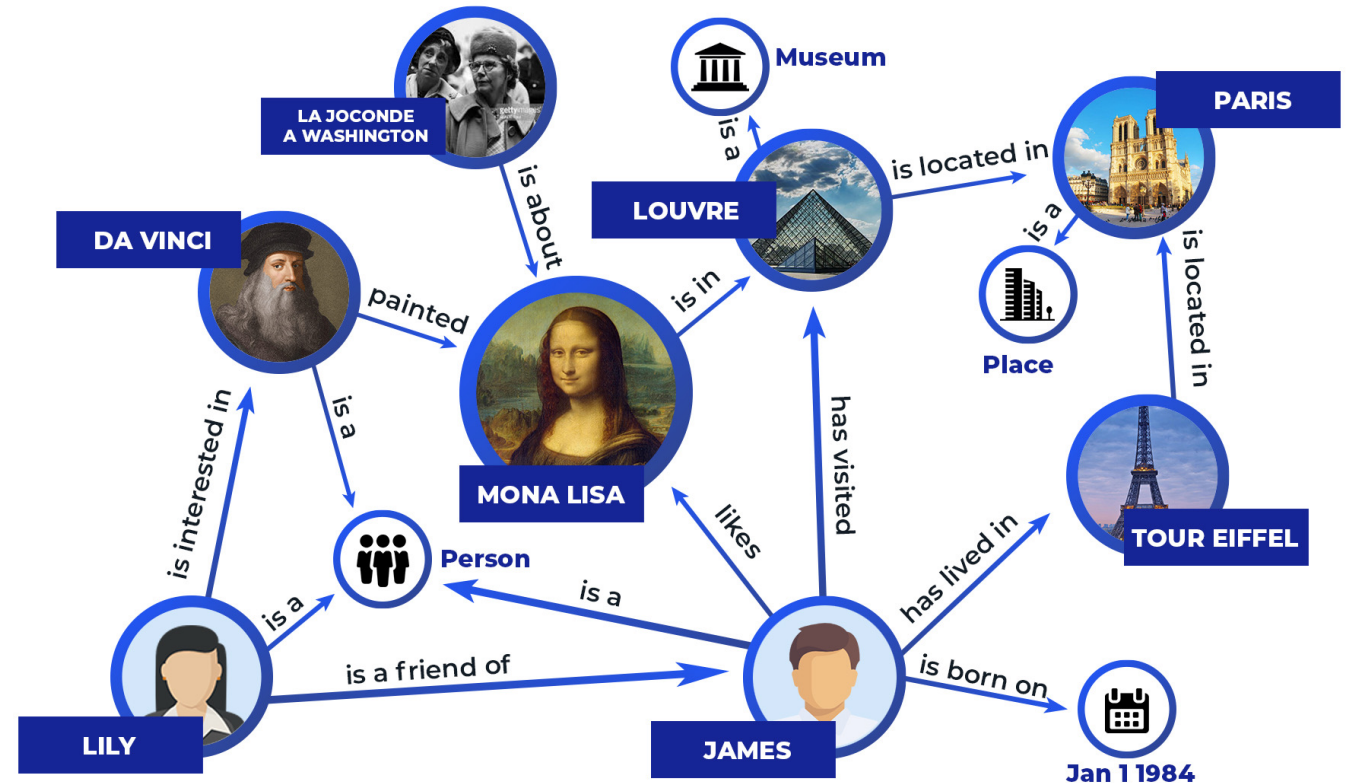
- Nodes: Classes/Objects
- Edges: Properties

Data Model: Resource Description Framework (RDF)

- Subject-Predicate-Object
- :JAMES :isBornOn :Jan-1-1984

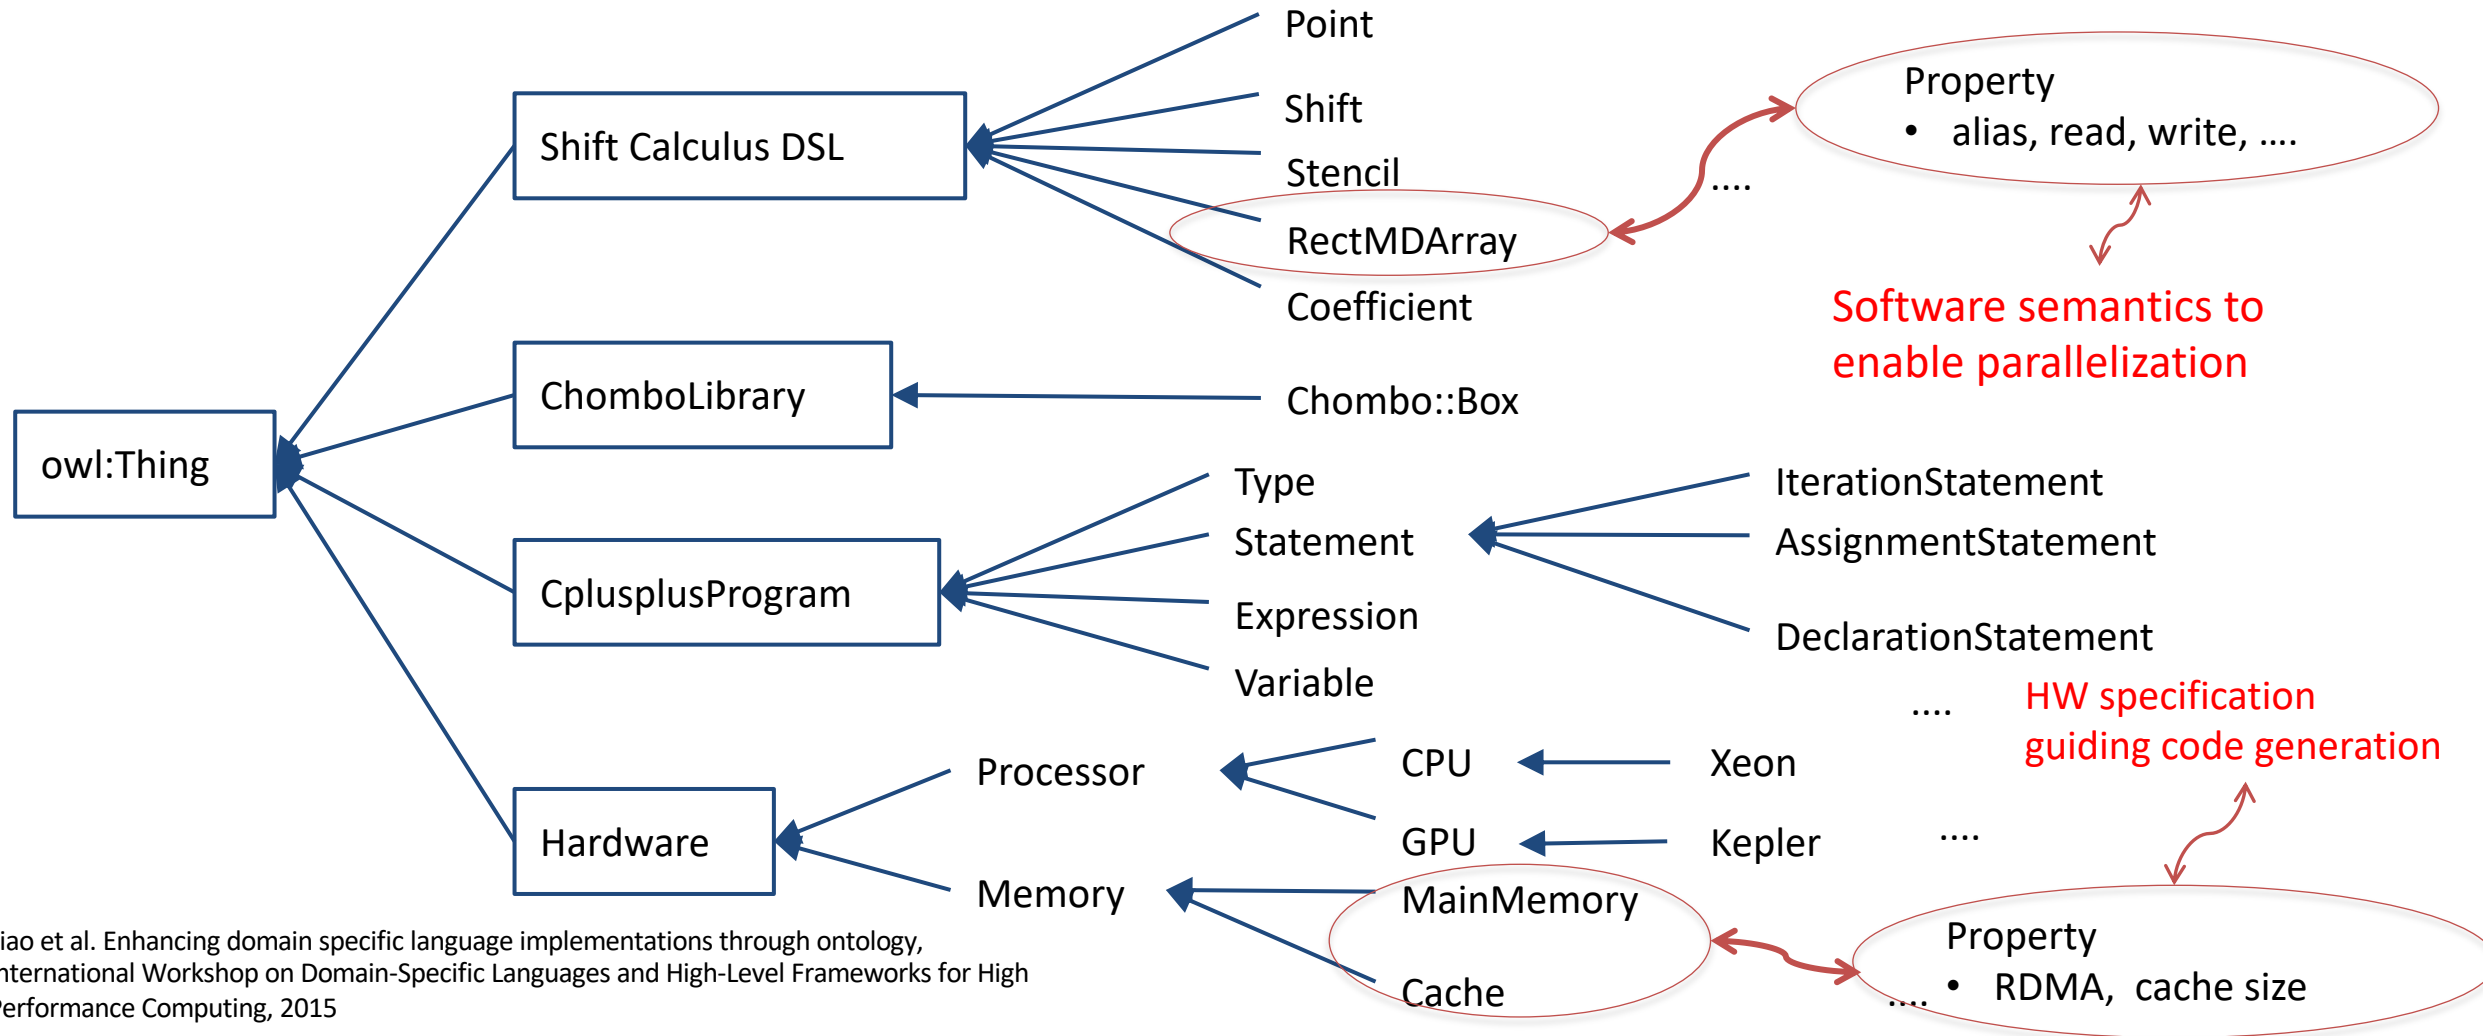
Queried using SPARQL

```
SELECT ?capital ?country
WHERE
{
  ?x ex:cityname ?capital ;
    ex:isCapitalOf ?y .
  ?y ex:countryname ?country ;
    ex:isInContinent ex:Africa .
}
```



Credit: <https://yashuseth.blog/2019/10/08/introduction-question-answering-knowledge-graphs-kgqa/>

# Building Vocabulary and Ontology for HPC




# Benefits of a Common Vocabulary and Knowledge Base

- Synergy between programming models:
  - Share/reuse/combine
    - domain-specific semantics
    - program analysis information
    - optimization rules
    - machine descriptions
  - Connect domain knowledge to programming models
  - Query semantics using standard query languages

- Some relevant studies

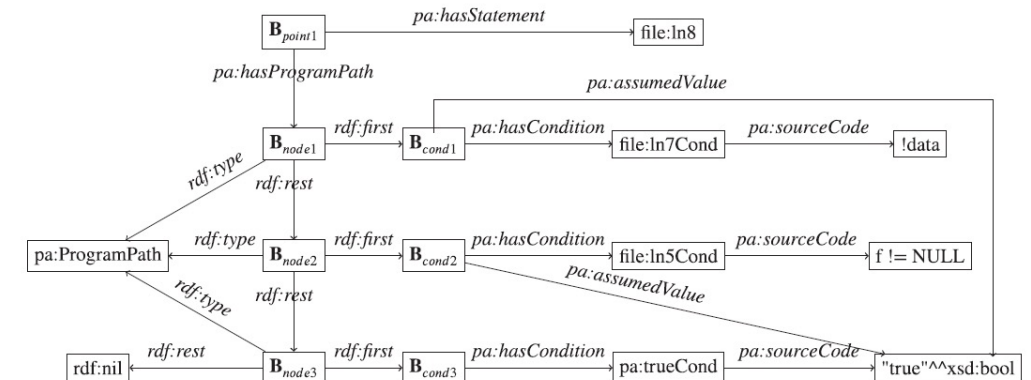
- Zhao et al. PATO: Towards Ontology-Based Program Analysis, The European Conference on Object-Oriented Programming, 2016
- Atzeni et al. CodeOntology: Querying Source Code in a Semantic Framework. In International Semantic Web Conference, 2017
- Pattipati et al. OPAL: An Extensible Framework for Ontology-based Program Analysis, Software Practice and Experience 2020

C program ontology for building mergeable, declarative analysis: PATO'16



Connecting parameters to semantics: CodeOntology'17

```
SELECT ? method
WHERE {
  ? method a : Method ;
  : hasParameter /: hasType : Double ;
  : associatedWith : Cube_root .
}
```



Path-sensitive analysis using semantics in a knowledge graph: OPAL'20

# Why Now?

## Maturing techniques and tool chains

- Web Ontology Language
- Resource Description Framework
- Large-scale knowledge graph databases

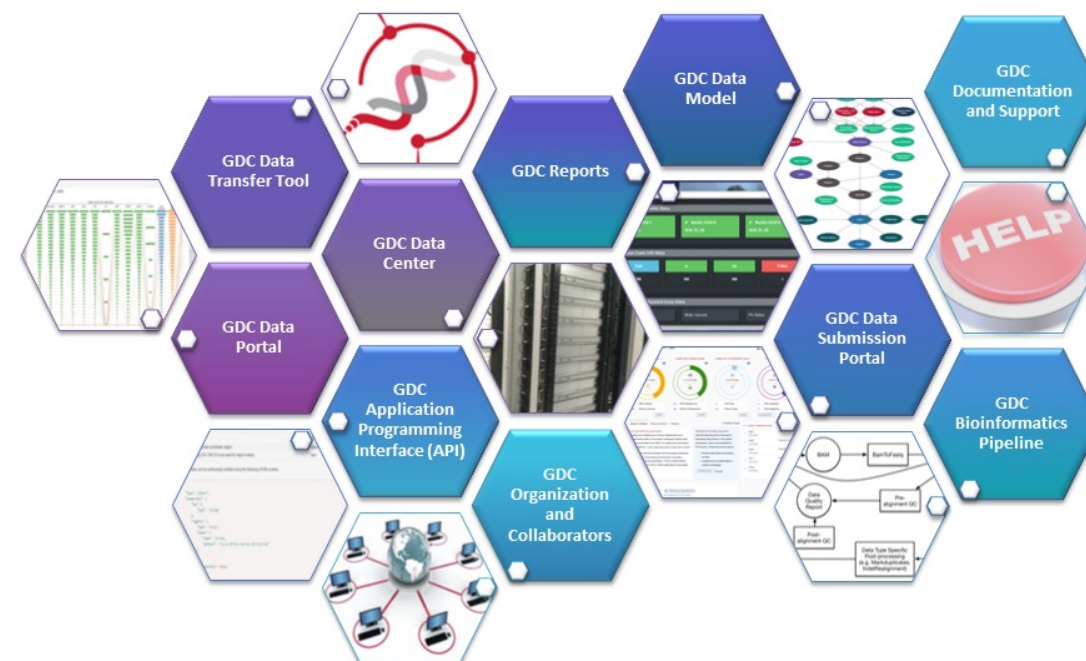
## Notable Vocabularies/Ontologies

- Schema.org, Wikidata, Yago 4, Linked Open Vocabularies (LOV)

## Promising prior research in ontologies + DSLs/compiler

## Notable Applications

- National Cancer Institute: Genomic Data Commons (GDC)



<https://gdc.cancer.gov/>

Linking knowledge from two domains:  
Clinical + Genomic data → Precision Cancer Treatments

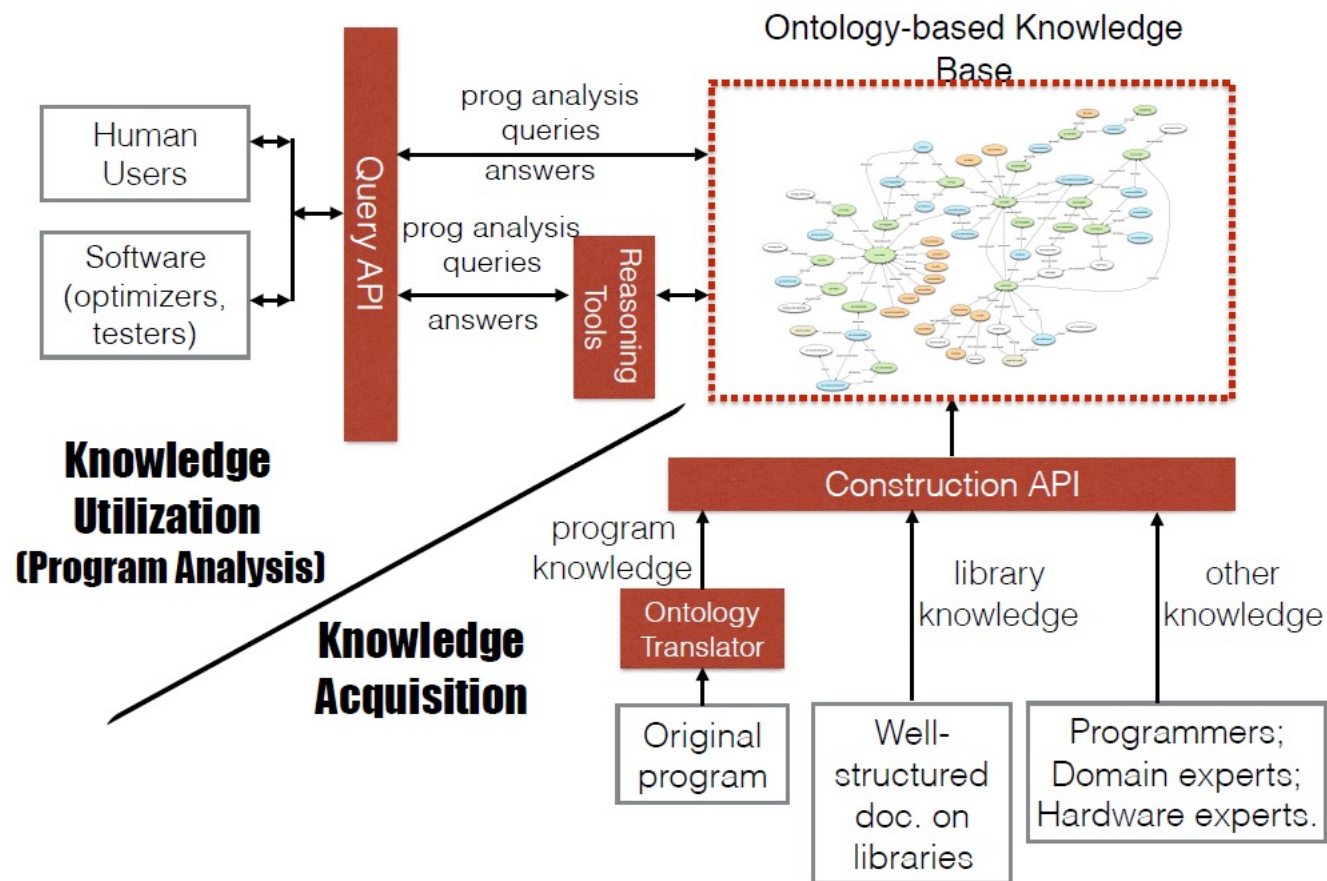
1



# Summary: CoDesign of Programming Models, powered by Common Vocabularies and Ontologies

HPC Programming models have a huge semantics gap

- **Need:** Heavily rely on sufficient knowledge of all layers of Software/hardware Stack for performance optimizations
- **Situation:** Mostly implicit or ad-hoc management of semantics/knowledge
- **Solution:** Formally manage knowledge across communities
  - Many existing knowledge management techniques available
  - Knowledge engineering: a community effort
  - The right time to address it is Now.



From PATO'16

# Disclaimer

---

This document was prepared as an account of work sponsored by an agency of the United States government. Neither the United States government nor Lawrence Livermore National Security, LLC, nor any of their employees makes any warranty, expressed or implied, or assumes any legal liability or responsibility for the accuracy, completeness, or usefulness of any information, apparatus, product, or process disclosed, or represents that its use would not infringe privately owned rights. Reference herein to any specific commercial product, process, or service by trade name, trademark, manufacturer, or otherwise does not necessarily constitute or imply its endorsement, recommendation, or favoring by the United States government or Lawrence Livermore National Security, LLC. The views and opinions of authors expressed herein do not necessarily state or reflect those of the United States government or Lawrence Livermore National Security, LLC, and shall not be used for advertising or product endorsement purposes.



# CASC

Center for Applied  
Scientific Computing



Sponsored by DOE Office of Science and LLNL.

#### Disclaimer

This document was prepared as an account of work sponsored by an agency of the United States government. Neither the United States government nor Lawrence Livermore National Security, LLC, nor any of their employees makes any warranty, expressed or implied, or assumes any legal liability or responsibility for the accuracy, completeness, or usefulness of any information, apparatus, product, or process disclosed, or represents that its use would not infringe privately owned rights. Reference herein to any specific commercial product, process, or service by trade name, trademark, manufacturer, or otherwise does not necessarily constitute or imply its endorsement, recommendation, or favoring by the United States government or Lawrence Livermore National Security, LLC. The views and opinions of authors expressed herein do not necessarily state or reflect those of the United States government or Lawrence Livermore National Security, LLC, and shall not be used for advertising or product endorsement purposes.

# Knowledge: hard to recover, critical to performance

- High-level abstractions in Source Codes written in C++
  - Encouraging code reuse and hiding implementation from interfaces to reduce software complexity
  - Functions, data structures, classes and templates ...
  - Could be standard or user-defined/domain-specific
- Semantics of high-level abstractions
  - Critical to optimizations, including parallelization
    - Read-only, data dependence, aliasing, accessed-by, etc.
    - Enabling a wide range of optimizations aimed for improving parallelism and reducing data movement.
  - Hard to be extracted by static/dynamic analysis
    - STL vector implementation ?→? elements are contiguous in memory
- Conventional compilers lose track of abstractions
  - Analyses and optimizations are mostly done on top of middle or low-level IR
  - Hard to trace back to the high-level abstractions represented in source level